

Document Retrieval with Embeddings

Joshua Cook

Introduction

Document
Retrieval with
Embeddings

Joshua Cook

Importance of Document Retrieval

- Powerful, cheap strategy for building new product features

Embedding-based Retrieval at Scribd



Author

Div Dasani

Published

April 12, 2021

Team

Recommendations

Tags:

machinelearning, real-
time, search, featured

Building recommendations systems like those implemented at large
companies like [Facebook](#) and [Pinterest](#) can be accomplished using off

<https://tech.scribd.com/blog/2021/embedding-based-retrieval-scribd.html>

How LinkedIn Is Using Embeddings to Up Its Match Game for Job Seekers



Jacob Mannix October 5, 2023



Co-Authors: [Jacob Mannix](#) and [Shaobo Zhang](#)

Think of how many times a day you use some type of search functionality across your devices and applications to discover information, find a contact, or a new job opportunity. The truth is we all depend on the ability to search for things online, and finding the right match to the information, organization, or to a job that maps to your skills and interests makes all the difference in our experiences and the knowledge we can gain.

The magic in search happens because of the technology that powers it. And one of the key technologies behind LinkedIn's search and recommendation features is embedding based retrieval (EBR). EBR helps deliver more relevant matches for our members and customers every day.

<https://engineering.linkedin.com/blog/2023/how-linkedin-is-using-embeddings-to-up-its-match-game-for-job-se>

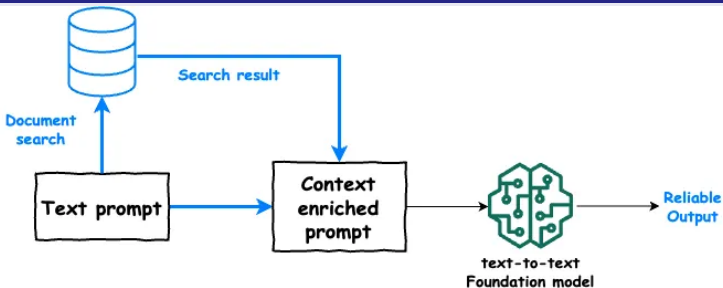
Traditional Methods and Their Limitations

- **Keyword-based Searching:** Explicitly matching terms
- **Boolean Queries:** Using logical operators (AND, OR, NOT)

Limitations:

- Lack of semantic understanding
- Low relevance
- Manual optimization needed

Retrieval-Augmented Generation



<https://learn.microsoft.com/en-us/azure/search/retrieval-augmented-generation-overview>


The Basics of Embeddings

Document
Retrieval with
Embeddings

Joshua Cook

What are Embeddings?

- Converting words to numerical space
- Example: Man \leftrightarrow Woman is similar to King \leftrightarrow Queen



Difference from Traditional Methods

- Understands semantic context
- Example: Differentiates “Apple” the company from “apple” the fruit

Demo

Document
Retrieval with
Embeddings

Joshua Cook

Brainstorming: Textual Datasets

Document
Retrieval with
Embeddings

Joshua Cook

(Three minutes, Groups of three)

Think of as many different textual datasets as you can three minutes.

Brainstorming: Building Context

Document
Retrieval with
Embeddings

Joshua Cook

(Five minutes, Groups of three)

Pick one set.

- Think of a context you can build with that data
 - What field or area or discipline or topic can this data be used to support GPT in answering questions?
- Name that context.
- Come up with five questions that you think GPT could answer give that context.

Brainstorming: Present

Document
Retrieval with
Embeddings

Joshua Cook

Share your context and questions with the cohort.